

Survival-guided matrix factorization identifies a reproducible tumor–stroma prognostic program in pancreatic cancer

Amber M. Young ^a, Alisa Yurovsky ^b, Xianlu Laura Peng ^{c,d}, Jen Jen Yeh ^{c,d,e}, Didong Li ^a, Naim U. Rashid ^{a,d,*},²

^aDepartment of Biostatistics, University of North Carolina at Chapel Hill, Chapel Hill, NC 27599, USA

^bDepartment of Biomedical Informatics, Stony Brook University, Stony Brook, NY 11794, USA

^cDepartment of Pharmacology, University of North Carolina at Chapel Hill, Chapel Hill, NC 27599, USA

^dLineberger Comprehensive Cancer Center, University of North Carolina at Chapel Hill, Chapel Hill, NC 27599, USA

^eDepartment of Surgery, University of North Carolina at Chapel Hill, Chapel Hill, NC 27599, USA

*Corresponding author: nur2@email.unc.edu

Abstract

In bulk tumor transcriptomes, the gene programs that explain the most variance are often not those that predict patient survival. We present DeSurv, a survival-supervised matrix factorization framework that jointly learns gene programs and their survival association, directing factorization toward outcome-relevant structure during discovery rather than post hoc. In pancreatic ductal adenocarcinoma (PDAC, TCGA + CPTAC), DeSurv recovered a compact three-program structure dominated by a co-varying tumor–stroma axis: Classical tumor with restCAF-like stroma at one pole, Basal-like tumor with proCAF-like stroma at the high-risk pole. Across five independent external cohorts, each standard-deviation shift toward the Basal-like/proCAF pole conferred a pooled hazard ratio of 1.50 (95% CI 1.31–1.72), with dichotomized high-risk patients showing HR 1.79 (95% CI 1.41–2.28); the score remained independently prognostic after adjustment for PurIST and DeCAF. By recovering programs that survive cross-cohort validation, DeSurv supports generalizable prognostic biomarker development from bulk transcriptomes.

Identifying the transcriptional programs that predict clinical outcomes is central to molecular subtyping and prognostic stratification in cancer¹. Pancreatic ductal adenocarcinoma (PDAC), among the deadliest solid tumors with a five-year relative survival rate of approximately 13%² and limited responsiveness to existing therapies, illustrates both the promise and the limits of this effort. Bulk transcriptomic studies have defined reproducible malignant programs, most consistently Basal-like and Classical tumor states^{1,3,4}, as well as stromal programs including tumor-promoting proCAF and tumor-restraining restCAF states⁵. Single-cell and spatial studies further show that PDAC tumors comprise malignant, stromal, immune and other microenvironmental components whose transcriptional states can co-vary within the same lesion^{6,7}. Whether a joint configuration of these programs carries prognostic information beyond discrete tumor and stromal labels, and whether such a configuration can be recovered de novo from bulk transcriptomes, remains incompletely understood.

Bulk RNA profiling remains the primary platform for cohort-scale survival annotation and prognostic modeling in large clinical cohorts^{8,9}, but each profile is a composite of malignant cells, cancer-associated fibroblasts, immune infiltrates and other cellular sources¹⁰. Nonnegative matrix factorization (NMF) has been instrumental for resolving this mixture into additive, interpretable gene programs^{3,11–13}. In PDAC, this strategy has yielded major biological insights, including molecular subtypes^{1,4,5,14}, virtual microdissection of Basal-like and Classical tumor programs³, and compartment-specific deconvolution of tumor and microenvironmental signals¹³.

A limitation of this workflow is that programs are usually discovered without outcome information and tested for clinical association only after factorization^{1,14}. Because unsupervised NMF minimizes reconstruction error, its factors tend to capture the highest-variance structure in the data, which may reflect tissue composition or other outcome-neutral variation rather than prognostic biology^{8,15}. In PDAC, exocrine and other low-purity signals can dominate unsupervised factors while carrying limited prognostic value^{4,5}. More broadly, tumor purity accounts for a large fraction of expression variation across cancer types¹⁶, illustrating how reconstruction-optimal and outcome-relevant subspaces can diverge. Reconstruction alone also leaves the factor basis non-identifiable^{17–19}, contributing to cross-study variability among unsupervised cancer subtype taxonomies^{20,21}.

This raises a biological and methodological question: which bulk-transcriptomic programs carry reproducible prognostic information in PDAC, and how are those programs distributed across malignant, stromal, immune and other components of the tumor microenvironment? We address this question through two non-circular tests: recovery of ground-truth program identities in simulations where the true gene sets are known a priori, and out-of-sample generalization to five independent PDAC cohorts on which no model parameters were estimated.

Here we present DeSurv, a survival-supervised deconvolution framework that integrates NMF with Cox proportional hazards modeling so that gene programs are shaped during discovery by both reconstruction and survival. Supervision is routed onto the gene-program matrix while preserving the mixture-coefficient interpretation of sample loadings, and trained programs score new samples by projection without retraining (Methods). In simulations, DeSurv recovers prognostic programs more reliably than standard NMF, with advantages that scale with the separation between high-variance and prognostic subspaces and vanish when no survival signal is present. In PDAC, DeSurv recovers a compact three-program structure dominated by a co-varying tumor–stroma axis linking Classical tumor identity and restCAF-like stroma with a Basal-like tumor and proCAF-like high-risk pole. This score generalizes across five independent PDAC cohorts without retraining and remains prognostic beyond established PurIST and DeCAF classifications. These findings identify tumor–stroma coupling as a reproducible prognostic program in PDAC and establish survival-supervised matrix factorization as a strategy for recovering outcome-relevant programs from bulk tumor transcriptomes wherever clinical outcomes are available.

Results

Model Overview

We developed DeSurv, a survival-supervised deconvolution framework that learns two linked quantities from bulk transcriptomes paired with patient outcomes: gene programs and their survival associations (Fig. 1A, Methods). DeSurv combines nonnegative matrix factorization with Cox proportional hazards modeling so that each gene program is shaped by both reconstruction of

the expression matrix and association with survival. The survival objective acts on the gene programs but not on the sample loadings, allowing the learned programs to be applied to new samples by projection without retraining. A supervision parameter controls the balance between reconstruction and survival, and both this parameter and the factorization rank are selected by Bayesian optimization on cross-validated concordance (Fig. 1B, Methods).

Applied to the PDAC training cohorts from TCGA²² and CPTAC²³ (Materials and Methods), cross-validated concordance varied across candidate ranks but plateaued from $k = 3$ through $k = 12$ (SI Appendix, Fig. S3), indicating that three supervised factors capture nearly as much prognostic information as twelve. Within this plateau, joint Bayesian optimization over rank and supervision strength with a one-standard-error rule (Methods) accordingly selected the parsimonious $k = 3$ at $\alpha = 0.334$. Peak cross-validated C-index was 0.655 across $n = 273$ patients with 139 events. Standard NMF rank-selection diagnostics (reconstruction residuals, cophenetic correlation, silhouette width)^{12,24} did not converge on a consistent rank in PDAC (SI Appendix, Fig. S4), motivating the supervised approach.

Survival supervision and standard NMF produce different factor structures in PDAC

To examine how survival supervision shapes the recovered factor structure, we fit standard NMF at the same rank as DeSurv ($k = 3$, selected above) and examined the overlap between each method’s factor-specific gene rankings and established PDAC gene programs (Fig. 2A–B). Matching k across methods isolates the effect of survival supervision on factor content from differences in model complexity; a sensitivity analysis at NMF’s independently selected ranks (elbow $k = 5$, BO $k = 7$) is provided in SI Appendix, Table S2. We first asked which malignant, stromal, immune and other microenvironmental programs were represented in each factorization, and then tested whether these structural differences corresponded to differences in survival contribution.

We refer to the three DeSurv factors as D1–D3 and the three NMF factors as N1–N3. In the DeSurv factorization, D1 correlated with Classical tumor programs and restCAF-like stromal signatures, D2 with proCAF-like programs, and D3 with Basal-like tumor programs (Fig. 2A). Among the programs detected, DeSurv organized the dominant survival-associated variation along a tumor–stroma axis, with D1 representing the Classical/restCAF-like low-risk pole and increasing risk reflecting movement toward Basal-like and proCAF-like structure, consistent with evidence that survival stratification in PDAC improves when tumor-intrinsic and CAF subtypes are considered jointly^{5,25}. No DeSurv factor showed a significant positive correlation with exocrine-associated gene programs (DECODER Exocrine, Collisson Exocrine, Bailey aberrantly differentiated endocrine–exocrine (ADEX); Fig. 2A). The survival contributions of these factors and their relationship to expression variance are quantified below.

Standard NMF produced a different factor organization (Fig. 2B). One factor (N2) was strongly correlated with exocrine-associated gene programs and negatively correlated with immune and stromal signatures (Fig. 2B). A second factor (N1) correlated with Classical tumor identity (Fig. 2B). A third factor (N3) correlated with immune, fibroblast, and extracellular matrix (ECM)-related programs (including Puleo Desmoplastic, SCISSORS proCAF and restCAF²⁶, Maurer Immune and ECM²⁷, and DECODER Activated Stroma), aggregating these into a single factor rather than separating them as DeSurv does (D1 and D2; Fig. 2A). At $k = 3$, NMF thus devotes one of its three factors to exocrine-associated variation, leaving two factors to capture all tumor-intrinsic and

microenvironmental programs. Whether these structural differences correspond to differences in prognostic information is quantified next.

We quantified per-factor contribution using two metrics (Fig. 2C; SI Appendix, Sections 10–11). For each factor, we computed the conditional variance explained (semi-partial R^2 : the incremental fraction of total variance accounted for by that factor given all other factors), and its survival contribution, measured as the change in Cox partial log-likelihood ($\Delta\ell$) when that factor is dropped from a k -factor Cox model. All three NMF factors had negligible $\Delta\ell$ despite spanning a range of variance, while DeSurv concentrated nearly all survival contribution into D1. Specifically, D1 explained 7% of variance with $\Delta\ell = 66.4$, while the highest-variance NMF factor explained 24.3% of expression variance but contributed only $\Delta\ell = 0.1$ to survival. The highest-variance DeSurv factor (D3, 35.4% of variance) also contributed minimally to survival after accounting for D1. Thus, in PDAC the highest-variance factors are not the most prognostic, consistent with tumor purity analyses across cancer types¹⁶.

To further characterize how the two factorizations relate, we examined pairwise W-matrix gene loading correlations restricted to the top 50 genes per factor (Fig. 2D). NMF’s tumor-identity factor (N1, top-correlated with Classical programs) was anti-correlated with DeSurv’s tumor factor (D3, top-correlated with Basal-like programs), reflecting genuine subtype discordance between Classical and Basal-like gene programs in the top-weighted genes of each method. NMF’s microenvironmental factor (N3) mapped primarily to DeSurv’s D2. Notably, DeSurv’s most prognostic factor (D1) showed only weak, roughly uniform correlations with all three NMF factors, indicating that the gene program structure underlying D1 does not have a direct counterpart in the unsupervised factorization. Because D1 couples Classical tumor and restCAF-like stromal signatures into a single survival-aligned program, its signal is distributed across multiple NMF factors, none of which individually captures the same tumor–stroma coupling. NMF’s exocrine-associated factor (N2) similarly lacked a strong DeSurv counterpart, with its strongest (though weak) correlation to D3, consistent with D3’s high variance and low survival contribution. Together, the two methods recover overlapping malignant and microenvironmental structure but differ in two key respects: NMF dedicates a factor to exocrine-compositional variation, whereas DeSurv does not allocate a dedicated factor to this low-prognostic component and instead constructs a survival-aligned tumor–stroma factor that NMF does not isolate. The net result is that DeSurv concentrates survival contribution into a single factor ($\Delta\ell = 66.4$), whereas no NMF factor achieves comparable survival contribution at the same rank. Whether this translates to improved generalization across independent cohorts is tested below.

Survival-aligned programs generalize across independent PDAC cohorts with lower complexity

A prognostic factorization is useful only if it generalizes beyond the training cohort. Having established that survival supervision and standard NMF produce different factor structures in the training data, we next tested whether each method’s structure transfers to independent cohorts without retraining. The five external validation cohorts (Dijk, Moffitt, PACA-AU array and RNA-seq, and Puleo) likewise consist of non-metastatic, treatment-naive specimens (SI Appendix, Table S1). We computed factor scores in each validation dataset by projecting the gene programs learned in the training cohort ($Z = \widetilde{W}^\top X_{\text{new}}$), yielding a continuous score for each factor per sample.

External validation C-index for NMF at independently selected ranks is summarized in SI Appendix, Table S2. Standard NMF at elbow-selected $k = 5$ and BO-selected $k = 7$ substantially improved

over NMF at $k = 3$, with $k = 7$ matching or exceeding DeSurv’s per-cohort C-index in most cohorts, though only when DeSurv’s tuning framework is applied ($\alpha = 0$ in the joint BO over rank and hyperparameters), since NMF’s own rank-selection heuristics (cophenetic, silhouette, residual; SI Appendix, Fig. S4) do not converge on this rank. However, this concordance gain still requires more than twice as many factors. DeSurv at $k = 3$ retains a significant adjusted hazard ratio after adjustment for PurIST and DeCAF (SI Appendix, Table S3), achieving independent prognostic value more parsimoniously. When β coefficients are re-fit on each validation cohort, DeSurv’s three-factor decomposition is BIC-preferred over NMF’s seven-factor decomposition in three of five cohorts, with the remaining two essentially tied (SI Appendix, Table S4). Across 5 independent PDAC cohorts ($n = 616$), the DeSurv linear predictor showed a consistent survival association (stratified Cox model: pooled HR per SD 1.50; 95% CI 1.31–1.72; $P = 2.6 \times 10^{-9}$; per-cohort forest plot in Fig. 3A; KM curves in SI Appendix, Fig. S5C–F). At the matched rank ($k = 3$), the analogous NMF linear predictor showed a weaker pooled effect (HR per SD 1.11; 95% CI 1.00–1.23; $P = 0.057$); a direct comparison using dichotomized risk groups is presented below.

To assess clinical interpretability of these continuous associations, we dichotomized validation samples into high- and low-risk groups using method-specific cutpoints, each independently optimized via cross-validation on the training data. The DeSurv linear predictor separated survival trajectories in the pooled validation cohort (HR = 1.79, 95% CI 1.41–2.28, $P < 0.001$; Fig. 3B). The same procedure applied to NMF at $k = 3$ yielded weaker stratification (HR = 1.38, 95% CI 0.92–2.09, $P = 0.122$; Fig. 3C). The DeSurv cutpoint classified a larger validation subset as high risk (151 vs 37 high-risk patients; 114 reclassified) while preserving strong survival separation. The DeSurv high-risk group was enriched for both established poor-prognosis subtypes (Basal-like by PurIST, proCAF by DeCAF; Fisher’s exact, SI Appendix, Fig. S6), capturing established consensus biology while adding independent prognostic information. The same analysis applied to NMF risk groups at $k = 3$ showed Basal-like enrichment but no significant proCAF enrichment (SI Appendix, Fig. S6), consistent with NMF’s failure to isolate the tumor–stroma coupling that underlies D1. Together, these results indicate that survival supervision yields a compact, biologically interpretable prognostic structure that generalizes across independent PDAC cohorts. Compared with standard NMF at the matched rank, DeSurv produced stronger validation survival stratification and recovered a tumor–stroma program enriched for both Basal-like tumor and proCAF-like stromal biology. Although higher-rank NMF models could achieve similar per-cohort concordance in some validation datasets, they required substantially greater complexity and did not yield the same parsimonious tumor–stroma coupling axis. DeSurv therefore identifies a lower-dimensional prognostic program that integrates established malignant and stromal PDAC subtype biology while retaining independent prognostic value beyond PurIST and DeCAF.

DeSurv recovers prognostic gene programs when variance and prognosis diverge

The cross-cohort generalization above is corroborated by simulations with known ground-truth latent structure. To test whether survival supervision improves gene program recovery under nonnegative constraints, we designed simulations in which prognostic gene programs explained low variance relative to outcome-neutral background signals, with survival times generated from latent factor scores via a Cox model ($p = 3,000$ genes, $n = 200$ samples, true $k = 3$; Materials and Methods, Simulation Studies). To test whether survival information should enter during factorization rather than only afterward, we compared DeSurv to standard NMF ($\alpha = 0$), where factors are learned by reconstruction error alone and survival associations are estimated by a separately fit Cox model. Both methods were tuned via BO over cross-validated C-index, so that differences in performance

reflect the role of supervision during factorization, not during model selection.

In the primary scenario, where prognostic programs explained low variance relative to outcome-neutral programs, DeSurv consistently achieved higher C-index (Fig. 4A) and substantially higher precision for recovering the true prognostic gene programs (Fig. 4B). Here precision denotes the fraction of a factor’s top-weighted genes (those with the largest factor specificity score s_{ij} , where n_{top} , the number of factor-specific genes entering the survival linear predictor, is chosen by BO; SI Appendix, Sections 6 and 9) that overlap with a true prognostic program’s gene set. The unsupervised baseline exhibited near-zero precision, indicating that variance-driven factorization failed to concentrate on the genes that actually drove survival. This confirms that outcome supervision during factorization is necessary to recover prognostic programs that do not dominate expression variance. Across 100 replicates, DeSurv achieved median C-index 0.837 versus 0.724 for standard NMF ($\Delta = 0.113$, paired Wilcoxon $P < 0.001$). The precision advantage exceeded the C-index advantage, reflecting that unsupervised factorization can match predictive performance while assigning different genes to each factor; survival supervision constrains gene-program identity itself. A similar decoupling appears empirically in PDAC at $k = 5$ (NMF’s elbow-selected rank), where supervised (DeSurv) and unsupervised (NMF) factorizations yield substantially different factor-specific top genes (SI Appendix, Fig. S7 [DeSurv $k = 5$] vs Fig. S8 [NMF $k = 5$]), with similar C-indices in external cohorts.

To verify that these gains reflect genuine signal recovery rather than overfitting, we repeated the analysis under a null scenario ($\beta = 0$) and a mixed scenario with partial overlap between variance-dominant and prognostic structure (SI Appendix, Fig. S2). DeSurv’s benefit scaled with the degree to which prognostic programs were masked by outcome-neutral variation: largest in the primary scenario ($\Delta = 0.113$), attenuated when variance and prognosis partially overlapped ($\Delta = 0.069$, paired Wilcoxon $P < 0.001$), and absent when no survival signal existed, where both methods yielded C-index near 0.5. DeSurv also recovered the true rank ($k = 3$) more frequently than standard NMF, which systematically under-selected near $k = 2$ (Fig. 4C).

Discussion

DeSurv is a survival-supervised matrix factorization framework that recovers prognostic transcriptional programs from bulk tumor data by coupling NMF reconstruction to a Cox survival objective during discovery; its architectural choice (supervision on W , mixture interpretation preserved on H) enables single-sample scoring by projection, a transportability property not shared by survival-supervised factorization methods that route supervision through H ^{28,29}. Applied to PDAC, this framework recovers a single prognostic axis that couples Classical tumor identity with restCAF stroma, distinct from tumor-intrinsic and stromal programs isolated separately. The recovered coupling generalizes across five independent PDAC cohorts without retraining and remains significant after adjustment for PurIST and DeCAF, suggesting that the prognostically relevant biology in PDAC is the joint configuration of tumor and stromal compartments rather than either alone.

In PDAC, the factor structure produced by survival supervision is biologically informative. The factor with the largest survival contribution (D1) couples Classical tumor programs with restCAF stromal signatures, consistent with growing evidence that clinical outcomes depend on the joint configuration of tumor-intrinsic and cancer-associated fibroblast subtypes^{5,25}. This co-occurrence within a single factor suggests that the prognostically relevant biology in PDAC is not purely tumor-intrinsic but reflects tumor–stroma coupling that unsupervised methods distribute across

multiple factors. Standard NMF, by contrast, allocated a dedicated factor to exocrine-compositional variation, a signal that reflects tumor purity rather than cancer biology⁴, leaving less capacity for the microenvironmental programs that DeSurv separates at the matched rank ($k = 3$). The contribution of this work is twofold. Methodologically, these programs are recovered de novo, without pre-specified signatures, with direct quantification of their survival associations during factorization. Biologically, this recovery surfaces tumor–stroma coupling as a single survival-aligned program — a relationship implicit in prior virtual microdissection³, experimental microdissection²⁷, and unsupervised deconvolution¹³ but not previously isolated as a single bulk-tx program.

Survival supervision concentrates prognostic signal into fewer factors, creating a broad concordance plateau that the one-standard-error rule leverages to select parsimonious ranks, a deliberate balance between biological completeness and clinical relevance. In PDAC, this plateau spanned $k = 3$ –12 (SI Appendix, Fig. S3); for standard NMF, concordance increased steadily with k , approaching DeSurv’s level by $k = 7$ (SI Appendix, Table S2) but requiring more than twice as many factors and producing a fragmented structure. Notably, NMF’s own unsupervised rank selection heuristics (reconstruction residuals, cophenetic correlation, and silhouette width; SI Appendix, Fig. S4) do not point to $k = 7$; the elbow criterion suggests $k = 5$, which does not match DeSurv’s concordance. NMF only reaches $k = 7$ through Bayesian optimization with $\alpha = 0$, that is, by applying DeSurv’s tuning infrastructure to an unsupervised model. Beyond this parsimony, DeSurv at $k = 3$ achieves independent prognostic value: both methods retain significance after classifier adjustment (SI Appendix, Table S3). The model selection framework (joint BO over rank and hyperparameters with the 1-SE rule) is itself part of the methodological contribution, not merely the survival gradient: parsimony is a property of the 1-SE rule, but supervision creates the flat concordance surface under which the rule can operate effectively.

Our simulations provide further guidance on when this tradeoff is most beneficial: DeSurv’s advantage is greatest when prognostic programs explain modest variance relative to outcome-neutral signals and vanishes under null conditions where no survival signal exists. DeSurv is therefore not intended to replace unsupervised NMF in all settings. When the goal is exploratory discovery without clinical endpoints, when survival data is sparse or unreliable, or when the application requires comprehensive biological characterization rather than prognostic stratification, unsupervised methods remain appropriate; recent reviews of cancer transcriptomic deconvolution methods provide broader guidance on method selection across these settings³⁰. The α parameter, tuned via Bayesian optimization, lets the data determine the appropriate balance for a given cancer type and cohort size, and the $\alpha = 0$ endpoint recovers standard NMF as a special case. DeSurv is likewise not proposed as a competing molecular classifier; translating gene programs into discrete clinical subtype labels would be a downstream application. The relevant comparison with existing classifiers is whether DeSurv’s factors capture independent prognostic signal beyond what those classifiers already explain; Table S3 confirms this directly, with the DeSurv $k = 3$ linear predictor retaining a significant adjusted hazard ratio after conditioning on PurIST and DeCAF (HR = 1.33, 95% CI 1.13–1.56, $P < 0.001$).

Several limitations define the scope of these findings. First, DeSurv assumes Cox proportional hazards, which may not hold in settings with delayed or non-proportional treatment effects; the convergence guarantee (SI Appendix, Theorem 1) applies to an idealized algorithm, and the implementation includes practical modifications that depart from the theoretical analysis, though empirically these do not affect convergence behavior. Extensions to more flexible survival models are a natural direction. Second, DeSurv was applied here only to PDAC, and broader benchmarking across cancer types is needed to establish the generality of the gap between reconstruction-optimal and

survival-optimal subspaces; the computational cost of Bayesian optimization may limit scalability to very large cohorts, though hyperparameter selection is performed once and the final model applies by projection. Third, all cohorts analyzed here consist of non-metastatic, treatment-naïve PDAC specimens. This homogeneity reduces confounding from treatment effects during factorization but limits direct applicability to treated patient populations where neoadjuvant chemotherapy is increasingly standard practice. Whether the gene programs identified here retain prognostic value in treated cohorts, or whether survival supervision in that setting recovers distinct treatment-responsive programs, remains an open question. Extending DeSurv to contemporary neoadjuvant trial cohorts (e.g., FOLFIRINOX and gemcitabine + nab-paclitaxel-treated series as they mature, where the projection-based scoring of the trained gene-program matrix can be applied without retraining) is therefore a natural and clinically important next step. Bulk profiling continues to underpin clinical molecular stratification, from trial programs such as COMPASS⁹ to CLIA-validated diagnostic assays in routine clinical care, a deployment landscape favored by the cost, FFPE compatibility, and standardization requirements of the clinical workflow. As single-cell cohorts with clinical annotation grow, DeSurv’s gene programs, which align with cell-type signatures confirmed by independent single-cell studies^{6,7}, can serve as a bridge between bulk-derived prognostic models and emerging cellular-resolution datasets.

The standard workflow in cancer transcriptomics (unsupervised discovery followed by retrospective evaluation) incurs a structural cost whenever the highest-variance signals are not the most prognostic, likely the norm in cancer types where tumor purity and tissue composition dominate expression variation¹⁶. DeSurv addresses this by aligning factorization with the clinical question from the outset: in PDAC, this recovered a Classical-tumor + restCAF program that transferred to five independent cohorts (pooled HR per SD 1.50) at a parsimonious $k = 3$ that survives PurIST/DeCAF adjustment. The supervised objective additionally pins the gene-program basis, addressing the basis-non-identifiability of reconstruction-only NMF: rotations of the factor matrix paired with compensating Cox coefficients can preserve overall predictive performance while shifting which genes load most strongly on each factor, so unsupervised gene programs need not reproduce across heterogeneous cohorts. This non-identifiability contributes to the cross-study variability observed among prior unsupervised PDAC subtype taxonomies^{1,3,14}. The combination of a survival-supervised gradient and joint Bayesian optimization with the one-standard-error rule provides a principled framework for extracting outcome-relevant programs from bulk tumor transcriptomes wherever clinical outcomes are available; as clinically annotated cohorts accumulate, this setting is more likely to expand alongside than be displaced by single-cell technologies.

Methods

Problem formulation and notation

Let $X \in \mathbb{R}_{\geq 0}^{p \times n}$ denote the nonnegative gene expression matrix. DeSurv approximates $X \approx WH$, where $W \in \mathbb{R}_{\geq 0}^{p \times k}$ contains nonnegative gene programs and $H \in \mathbb{R}_{\geq 0}^{k \times n}$ contains sample loadings. Additionally, let $y \in \mathbb{R}_{\geq 0}^n$ denote patient survival times, $\delta \in \{0, 1\}^n$ the censoring indicators ($\delta_i = 1$ if the event is observed), and $\beta \in \mathbb{R}^k$ the Cox regression coefficients. Survival outcomes are modeled through a Cox proportional hazards model with factor scores $Z = W^\top X \in \mathbb{R}^{k \times n}$ and linear predictor $Z^\top \beta$.

The DeSurv Model

DeSurv integrates NMF with penalized Cox regression to identify gene programs associated with survival outcomes. The method operates in a semi-supervised framework: the reconstruction loss penalizes deviation from the observed expression data, while the survival term directs gene programs toward outcome-relevant structure. The parameter $\alpha \in [0, 1)$ controls the relative contribution of each objective.

The joint objective is

$$\mathcal{L}(W, H, \beta) = (1 - \alpha) \mathcal{L}_{\text{NMF}}(W, H) - \alpha \mathcal{L}_{\text{Cox}}(W, \beta), \quad (1)$$

where $\mathcal{L}_{\text{NMF}}(W, H)$ is the NMF reconstruction error (including an ℓ_2 penalty on H) and $\mathcal{L}_{\text{Cox}}(W, \beta)$ is the elastic-net penalized partial log-likelihood. A critical architectural choice is where survival supervision enters the factorization. Because factor scores are defined as $Z = W^\top X$, the Cox partial likelihood $\mathcal{L}_{\text{Cox}}$ is a direct function of W and β , and the survival gradient acts explicitly on the gene program matrix W . The sample-level loadings H enter only through the reconstruction term and receive no survival gradient, preserving their interpretation as mixture coefficients^{11,31}. When $\alpha = 0$, DeSurv reduces to standard unsupervised NMF. The Cox component can also accommodate additional sample-level covariates, such as tumor stage or grade, during optimization.

Optimization alternates updates for H , W , and β (Algorithm 1). The reconstruction term is minimized using multiplicative updates that preserve nonnegativity; the supervision term enters W 's gradient as a Cox partial-likelihood derivative; and β is updated by Coxnet coordinate descent on the current factor scores $Z = W^\top X$ with elastic-net penalty parameters (λ, ξ) . Although the joint problem is non-convex, the alternating updates converge to a stationary point under mild regularity conditions (SI Appendix, Theorem 1). Closed-form update equations, the convergence proof, and BO control parameters are given in SI Appendix, Sections 1–3.

Algorithm 1 DeSurv block-coordinate descent

Input: Expression matrix $X \in \mathbb{R}_{\geq 0}^{p \times n}$, survival (y, δ) , supervision strength $\alpha \in [0, 1)$, rank k , regularization (λ, ξ) , tolerance ϵ

Output: $W \in \mathbb{R}_{\geq 0}^{p \times k}$, $H \in \mathbb{R}_{\geq 0}^{k \times n}$, $\beta \in \mathbb{R}^k$

- 1: Initialize $W^{(0)}$ via consensus-based seeding (SI Appendix, Section 6); set $H^{(0)}, \beta^{(0)}$
 - 2: **repeat**
 - 3: $H \leftarrow$ multiplicative-update minimizer of $\mathcal{L}_{\text{NMF}}(W, \cdot)$
 - 4: $W \leftarrow$ multiplicative update with combined gradient $\nabla \mathcal{L} = (1 - \alpha) \nabla_W \mathcal{L}_{\text{NMF}} - \alpha \nabla_W \mathcal{L}_{\text{Cox}}(W, \beta)$
 - 5: $\beta \leftarrow$ Coxnet coordinate descent on $Z = W^\top X$ with penalty (λ, ξ)
 - 6: **until** $|\mathcal{L}^{(t)} - \mathcal{L}^{(t-1)}| / |\mathcal{L}^{(t-1)}| < \epsilon$
-

Hyperparameter selection and cross-validation

Hyperparameters $(k, \alpha, \lambda, \xi, n_{\text{top}})$ were selected by maximizing the cross-validated C-index using Bayesian optimization (BO)^{32,33}, which finds near-optimal configurations using far fewer objective evaluations than grid or random search over mixed integer–continuous hyperparameter spaces, with final rank k chosen by the one-standard-error rule (the smallest k whose predicted performance lies within one standard error of the maximum). All model selection was performed entirely within

training data using nested cross-validation; no validation cohort data were used during any stage of tuning. Each fold was trained using multiple random initializations, and fold-level performance was defined as the average C-index across initializations. For stability, we used a consensus-based initialization for the final model, aggregating multiple DeSurv runs into a gene-gene co-occurrence matrix and constructing an initialization W_0 from the resulting clusters (SI Appendix, Section 6). Before validation, each column of W was truncated to its BO-selected number of top genes (details in SI Appendix, Section 6), denoted $\widetilde{W}$. Columns of W were normalized using the same convention applied during training (column ℓ_2 -norm scaling; SI Appendix, Section 6), and the resulting truncated matrix $\widetilde{W}$ was held fixed for all validation cohorts. In PDAC, BO selected $n_{\text{top}} = 270$ top genes per factor.

Because the goal is prognostic factorization, DeSurv selects k using the criterion that directly measures prognostic performance (cross-validated C-index) rather than the unsupervised diagnostics commonly used for NMF rank selection^{12,24}. In the PDAC training data, standard NMF diagnostics (reconstruction residuals, cophenetic correlation, mean silhouette width) yielded inconsistent guidance across candidate ranks, a known limitation of unsupervised rank-selection criteria that can disagree on the same dataset (SI Appendix, Fig. S4). Cross-validated concordance varied across candidate ranks but plateaued from $k = 3$ through $k = 12$ (SI Appendix, Fig. S3). The one-standard-error rule applied to joint BO over rank and supervision strength accordingly selected the parsimonious $k = 3$ at $\alpha = 0.334$.

External validation cohorts were scored by projecting new datasets onto the learned programs via $Z = \widetilde{W}^\top X_{\text{new}}$, and survival associations were evaluated using C-index, hazard ratios, and log-rank statistics. Because the survival-supervised information lives in W rather than H , scoring is a pure matrix projection: no retraining, no sample-level nonnegative least-squares estimation on the new data, and no validation survival information is required at the scoring step. This is a property not shared by methods that route survival supervision through H , which require sample-loading re-estimation on each new cohort^{28,29}. Validation survival outcomes were therefore used only for downstream performance evaluation, not for model training, scoring, or hyperparameter tuning. Reported external hazard ratios and log-rank statistics therefore reflect purely out-of-sample generalization. Further training and validation details appear in SI Appendix, Part I.

Simulation studies

Simulation studies were conducted to assess recovery of prognostic latent structure and survival prediction under controlled conditions. Gene expression data were generated from a nonnegative factor model $X = WH$, where gene loadings W comprised three gene classes: marker genes, background genes, and noise genes. Marker genes were simulated to load strongly on a single factor and weakly on others, background genes to load strongly across all factors, and noise genes to have uniformly low loadings. Specifically, marker gene primary loadings were drawn from Gamma(shape = 3, rate = 0.8) and cross-loadings onto other factors from Gamma(1, 20); background gene loadings from Gamma(2, 1) across all factors; and noise gene loadings from Gamma(1, 20). Sample-level factor activities H were drawn from Gamma(2, 2). Each dataset comprised $G = 3,000$ genes, $N = 200$ samples, and $K = 3$ factors, with 150 marker genes per factor and 500 shared background genes. Cox regression coefficients were $\beta = (2, 0, 0)^\top$ in the primary scenario, $\beta = 0$ in the null scenario, and $\beta = (2, 0, 0)^\top$ in the mixed scenario where survival risk was driven by 300 genes comprising the 150 factor-1 markers and 150 randomly sampled background genes, creating partial overlap between the prognostic signal and variance-dominant outcome-neutral

programs.

Survival times were generated from an exponential distribution in which risk depended on marker gene expression through $X^T W_{\text{true}}$, where W_{true} retained marker gene loadings for their corresponding factors and was zero otherwise; censoring times were generated independently from an exponential distribution. Each dataset was analyzed using DeSurv and standard NMF followed by Cox regression on inferred factors, both tuned using cross-validated concordance index via Bayesian optimization. Performance was summarized across repeated simulation replicates. We considered three simulation scenarios: a primary scenario where prognostic programs explain low variance relative to outcome-neutral signals, a null scenario with no survival signal, and a mixed scenario where prognostic and variance-dominant programs partially overlap.

Real-world datasets

We analyzed publicly available RNA sequencing (RNA-seq) and microarray cohorts of PDAC with corresponding overall survival outcomes. Gene expression matrices were analyzed on a log scale: RNA-seq cohorts as $\log_2(\text{TPM} + 1)$; microarray cohorts in platform-normalized log-scale intensities. For each training cohort separately, genes were ranked by mean expression and by variance independently; the top 3,000 genes with the lowest sum of these two ranks were retained, selecting genes that are both highly expressed and highly variable; the intersection across training cohorts yielded 1,970 genes used for all analyses. Validation datasets were restricted to this same gene set. Survival times and censoring indicators were taken from the associated clinical annotations. Of the seven PDAC cohorts we considered, two were used for training (TCGA and CPTAC) and five were used for external validation (Dijk, Moffitt, PACA-AU array, PACA-AU RNA-seq, and Puleo). To harmonize differences in scale across cohorts, filtered gene expression data were within-subject rank transformed before model training (ranks are inherently nonnegative, preserving NMF compatibility). More details about the datasets can be found in SI Appendix, Section 7.

Per-cohort participant flow is summarized in SI Appendix, Table S1. Of 321 pooled training-cohort samples (TCGA-PAAD, $n = 181$; CPTAC-3, $n = 140$), 48 were excluded for missing survival time, missing event indicator, or quality-control failure, yielding 273 analytic training samples (139 events). Of 979 pooled validation-cohort samples, 363 were excluded for non-PDAC histology, non-tumor tissue, or missing survival annotation, yielding 616 analytic validation samples (414 events) across five cohorts: Dijk ($n = 90$, 81 events), Moffitt GEO array ($n = 123$, 83), PACA-AU array ($n = 63$, 38), PACA-AU RNA-seq ($n = 52$, 31), and Puleo array ($n = 288$, 181). With 139 events and 3 factor-score covariates (one per DeSurv factor at $k = 3$), the Cox-stage events-per-variable ratio is 46, well above the conventional 10-EPV threshold, supporting estimation stability.

Data, Materials, and Software Availability

All datasets used in this study are publicly available and de-identified; no Institutional Review Board (IRB) approval was required. Training data were obtained from the Genomic Data Commons: TCGA-PAAD³⁴ and CPTAC-3³⁵. Validation cohorts were obtained from ArrayExpress (Dijk, E-MTAB-6830³⁶; Puleo, E-MTAB-6134³⁷), the Gene Expression Omnibus (GEO; Moffitt, GSE71729³⁸), and the International Cancer Genome Consortium (ICGC) Data Portal (PACA-AU array and RNA-seq; European Genome-phenome Archive (EGA) study EGAS00001000154³⁹). Cohort details, sample sizes after quality control, and patient-level metadata used are summarized in SI Appendix,

Section 7. Expression data and molecular classifications (PurIST, DeCAF) for the validation cohorts were originally compiled for the DeCAF study⁵.

All analyses were performed in R (version 4.6.0; R Core Team) using the following key packages: DeSurv (v1.0.1; this work), NMF³¹, glmnet⁴⁰, survival, ggplot2, and cowplot. Bayesian optimization for hyperparameter selection was implemented within DeSurv. An R package implementing DeSurv is available at github.com/rashidlab/DeSurv (archived at Zenodo: DOI 10.5281/zenodo.20150929). Code, processed data, and the exact session-info file (R + package versions used to render this manuscript) are available at github.com/rashidlab/DeSurv-paper (archived at Zenodo: DOI 10.5281/zenodo.20150946).

Data and Code Availability

The DeSurv R package is available at <https://github.com/rashidlab/DeSurv> (archived at Zenodo, DOI 10.5281/zenodo.20150929). Reproducibility code, processed data, and session-info for this manuscript are available at <https://github.com/rashidlab/DeSurv-paper> (archived at Zenodo, DOI 10.5281/zenodo.20150946). All training and validation cohorts are publicly available: TCGA-PAAD and CPTAC-3 from the Genomic Data Commons; Dijk (E-MTAB-6830) and Puleo (E-MTAB-6134) from ArrayExpress; Moffitt (GSE71729) from the Gene Expression Omnibus; and PACA-AU (EGA study EGAS00001000154) from the ICGC Data Portal.

Acknowledgements

This work was supported by the National Cancer Institute grants U01 CA274298, P50 CA257911, T32 CA106209, and R01 CA288145, and by the Department of Defense Peer Reviewed Cancer Research Program (HT9425241103100121). We thank the patients and research teams who contributed samples and clinical data to the public cohorts used here (TCGA, CPTAC, Dijk, Moffitt, Puleo, PACA-AU).

Inclusion and Diversity. The patient cohorts analyzed in this study were curated by the original investigators or consortia (TCGA, CPTAC, ICGC, and the original Moffitt, Dijk, and Puleo studies) and our access was limited to the de-identified expression and survival data they provide; we did not stratify analyses by sex, race, or ethnicity because consistent demographic annotations were not available across all cohorts. Sex and other demographic variables that could modify the prognostic relevance of the gene programs identified here remain important directions for follow-up work as more comprehensively annotated cohorts become available.

Author Contributions

Conceptualization: A.M.Y., N.U.R. Methodology: A.M.Y., A.Y., D.L., N.U.R. Software: A.M.Y. Formal analysis: A.M.Y. Investigation: A.M.Y. Data curation: A.M.Y., X.L.P. Visualization: A.M.Y. Writing – original draft: A.M.Y., N.U.R. Writing – review and editing: A.M.Y., A.Y., X.L.P., J.J.Y., D.L., N.U.R. Resources (PDAC domain expertise and biological interpretation): X.L.P., J.J.Y. Supervision: N.U.R. Funding acquisition: A.M.Y., N.U.R. All authors reviewed and approved the final manuscript.

Competing Interests

N.U.R. and J.J.Y. are named as inventors on US patents US20220127676A1 and US20250066856A1, which relate to pancreatic cancer subtype classification (establishment of Basal and Classical subtypes and their prediction) derived from prior work and were not used directly in this study. The other authors declare no competing interests.


References

1. Collisson, E. A. *et al.* Subtypes of pancreatic ductal adenocarcinoma and their differing responses to therapy. *Nature Medicine* **17**, 500–503 (2011).
2. Siegel, R. L., Kratzer, T. B., Wagle, N. S., Sung, H. & Jemal, A. Cancer statistics, 2026. *CA: A Cancer Journal for Clinicians* <https://doi.org/10.3322/caac.70043> (2026) doi:10.3322/caac.70043.
3. Moffitt, R. A. *et al.* Virtual microdissection identifies distinct tumor- and stroma-specific subtypes of pancreatic ductal adenocarcinoma. *Nature Genetics* **47**, 1168–1178 (2015).
4. Rashid, N. U. *et al.* Purity independent subtyping of tumors (PurIST), a clinically robust, single-sample classifier for tumor subtyping in pancreatic cancer. *Clinical Cancer Research* **26**, 82–92 (2020).
5. Peng, X. L. *et al.* DeCAF defines clinical fibroblast subtypes and multidimensional tumor-stroma crosstalk shaping prognosis and immunotherapy response. *Cell Reports Medicine* **7**, 102611 (2026).
6. Chan-Seng-Yue, M. *et al.* Transcription phenotypes of pancreatic cancer are driven by genomic events during tumor evolution. *Nature Genetics* **52**, 231–240 (2020).
7. Werba, G. *et al.* Single-cell RNA sequencing reveals the effects of chemotherapy on human pancreatic adenocarcinoma and its tumor microenvironment. *Nature Communications* **14**, 797 (2023).
8. Bair, E. & Tibshirani, R. Semi-supervised methods to predict patient survival from gene expression data. *PLoS Biology* **2**, e108 (2004).
9. Aung, K. L. *et al.* Genomics-driven precision medicine for advanced pancreatic cancer: Early results from the COMPASS trial. *Clinical Cancer Research* **24**, 1344–1354 (2018).
10. Nguyen, H., Nguyen, H., Tran, D., Draghici, S. & Nguyen, T. Fourteen years of cellular deconvolution: Methodology, applications, technical evaluation and outstanding challenges. *Nucleic Acids Research* **52**, 4761–4783 (2024).
11. Lee, D. D. & Seung, H. S. Learning the parts of objects by non-negative matrix factorization. *Nature* **401**, 788–791 (1999).
12. Brunet, J.-P., Tamayo, P., Golub, T. R. & Mesirov, J. P. Metagenes and molecular pattern discovery using matrix factorization. *Proceedings of the National Academy of Sciences USA* **101**, 4164–4169 (2004).
13. Peng, X. L., Moffitt, R. A., Torphy, R. J., Volmar, K. E. & Yeh, J. J. De novo compartment deconvolution and weight estimation of tumor samples using DECODER. *Nature Communications* **10**, 4729 (2019).

14. Bailey, P. *et al.* Genomic analyses identify molecular subtypes of pancreatic cancer. *Nature* **531**, 47–52 (2016).
15. Cook, R. D. Fisher lecture: Dimension reduction in regression. *Statistical Science* **22**, 1–26 (2007).
16. Aran, D., Sirota, M. & Butte, A. J. Systematic pan-cancer analysis of tumour purity. *Nature Communications* **6**, 9971 (2015).
17. Donoho, D. & Stodden, V. When does non-negative matrix factorization give a correct decomposition into parts? in *Advances in neural information processing systems 16 (NIPS 2003)* (eds. Thrun, S., Saul, L. K. & Schölkopf, B.) 1141–1148 (MIT Press, 2004).
18. Laurberg, H., Christensen, M. G., Plumbly, M. D., Hansen, L. K. & Jensen, S. H. Theorems on positive data: On the uniqueness of NMF. *Computational Intelligence and Neuroscience* **2008**, 764206 (2008).
19. Gillis, N. The why and how of nonnegative matrix factorization. in *Regularization, optimization, kernels, and support vector machines* (eds. Suykens, J. A. K., Signoretto, M. & Argyriou, A.) 257–291 (Chapman; Hall/CRC, 2014).
20. Lautizi, M., Baumbach, J., Weichert, W., Steiger, K. & List, M. A systematic assessment of pancreatic cancer molecular subtype classification across cohorts. *NAR Cancer* **4**, zcac030 (2022).
21. Torre-Healy, L. A. *et al.* Open-source curation of a pancreatic ductal adenocarcinoma gene expression analysis platform (pdacR) supports a two-subtype model of clinical relevance. *Communications Biology* **6**, 163 (2023).
22. Raphael, B. J. *et al.* Integrated genomic characterization of pancreatic ductal adenocarcinoma. *Cancer Cell* **32**, 185–203 (2017).
23. Ellis, M. *et al.* Connecting genomic alterations to cancer biology with proteomics: The NCI clinical proteomic tumor analysis consortium. *Cancer Discovery* **3**, 1108–1112 (2013).
24. Frigyesi, A. & Höglund, M. Non-negative matrix factorization for the analysis of complex gene expression data: Identification of clinically relevant tumor subtypes. *Cancer Informatics* **6**, 275–292 (2008).
25. Elyada, E. *et al.* Cross-species single-cell analysis of pancreatic ductal adenocarcinoma reveals antigen-presenting cancer-associated fibroblasts. *Cancer Discovery* **9**, 1102–1123 (2019).
26. Leary, J. R. *et al.* Sub-cluster identification through semi-supervised optimization of rare-cell silhouettes (SCISSORS) in single-cell RNA-sequencing. *Bioinformatics* **39**, btad449 (2023).
27. Maurer, C. *et al.* Experimental microdissection enables functional harmonisation of pancreatic cancer subtypes. *Gut* **68**, 1034–1043 (2019).
28. Huang, Z., Salama, P., Shao, W., Zhang, J. & Huang, K. Low-rank reorganization via proportional hazards non-negative matrix factorization unveils survival associated gene clusters. <https://arxiv.org/abs/2008.03776> (2020).
29. Le Goff, V. *et al.* SurvNMF: Non-negative matrix factorization supervised for survival data analysis. (2025).
30. Dai, Y. *et al.* A guide to transcriptomic deconvolution in cancer. *Nature Reviews Cancer* **26**, 84–103 (2025).
31. Gaujoux, R. & Seoighe, C. A flexible R package for nonnegative matrix factorization. *BMC Bioinformatics* **11**, 367 (2010).

32. Snoek, J., Larochelle, H. & Adams, R. P. Practical Bayesian optimization of machine learning algorithms. in *Advances in neural information processing systems 25 (NeurIPS 2012)* (2012).
33. Shahriari, B., Swersky, K., Wang, Z., Adams, R. P. & Freitas, N. de. Taking the human out of the loop: A review of Bayesian optimization. *Proceedings of the IEEE* **104**, 148–175 (2016).
34. The Cancer Genome Atlas Research Network. Data from ‘integrated genomic characterization of pancreatic ductal adenocarcinoma’ (TCGA-PAAD). (2017).
35. National Cancer Institute Clinical Proteomic Tumor Analysis Consortium. Data from CPTAC-3 pancreatic ductal adenocarcinoma cohort. (2019).
36. Dijk, F. *et al.* Data from ‘unsupervised class discovery in pancreatic ductal adenocarcinoma’. (2020).
37. Puleo, F. *et al.* Data from ‘stratification of pancreatic ductal adenocarcinomas based on tumor and microenvironment features’. (2018).
38. Moffitt, R. A. *et al.* Data from ‘virtual microdissection identifies distinct tumor- and stroma-specific subtypes of pancreatic ductal adenocarcinoma’. (2015).
39. Australian Pancreatic Cancer Genome Initiative & International Cancer Genome Consortium. Pancreatic cancer (PACA-AU) cohort: Array and RNA-sequencing data. (2016).
40. Simon, N., Friedman, J. H., Hastie, T. & Tibshirani, R. Regularization paths for cox’s proportional hazards model via coordinate descent. *Journal of Statistical Software* **39**, 1–13 (2011).

(A) The DeSurv model



(B) Data-driven model selection via Bayesian optimization




Figure 1: Overview of the DeSurv framework. (A) DeSurv jointly optimizes NMF reconstruction ($X \approx WH$) and Cox proportional hazards likelihood. The gene program matrix W is shared between objectives: factor scores $Z = W^T X$ serve as covariates in the Cox model with coefficients β . A supervision parameter $\alpha \in [0, 1]$ controls the balance between reconstruction fidelity ($\alpha = 0$, standard NMF) and survival prediction (α close to 1). (B) The factorization rank (k), supervision strength (α), and regularization (λ) are selected via Bayesian optimization over cross-validated concordance index.




Figure 2: Survival supervision produces different factor structures in PDAC (TCGA + CPTAC training cohort, $n = 273$ patients, 139 events). (A–B) Spearman correlations between factor gene rankings (top 50 genes per factor) and established PDAC gene programs for DeSurv (A) and standard NMF (B) at $k = 3$. Only $|r| > 0.2$ shown. DeSurv concentrates survival signal into D1 (Classical tumor + restCAF-like stroma) while suppressing exocrine variation; standard NMF devotes one factor to exocrine signals. (C) Fraction of expression variance explained versus survival contribution (Δ Cox partial log-likelihood upon factor removal) for each factor. DeSurv concentrates survival signal into D1 while the highest-variance NMF factor (N2, exocrine-associated) contributes negligibly to survival. (D) Pairwise correspondence between NMF and DeSurv factors, measured by Spearman rank correlation of W-matrix gene loadings restricted to the top 50 genes per factor.




Figure 3: Survival-aligned programs generalize across independent PDAC cohorts (pooled $n = 616$ patients across 5 cohorts: Dijk $n = 90$; Moffitt GEO array $n = 123$; PACA-AU array $n = 63$; PACA-AU RNA-seq $n = 52$; Puleo array $n = 288$). (A) Forest plot of univariate hazard ratios (95% CIs) from separate Cox models for each factor, showing each factor's marginal survival association in each validation cohort. DeSurv factor D1 shows a consistent protective association across all five cohorts (pooled HR = 0.76), while no standard NMF factor achieves a comparable pooled effect. Error bars are 95% confidence intervals; black diamonds denote inverse-variance-weighted pooled estimates. (B-C) Kaplan-Meier curves for pooled external validation samples stratified by the combined multivariate linear predictor ($\hat{\beta}_1 Z_1 + \hat{\beta}_2 Z_2 + \hat{\beta}_3 Z_3$, where coefficients are learned from training data), dichotomized at a cross-validated optimal cutpoint: DeSurv (B) and standard NMF (C).




Figure 4: Survival supervision improves recovery of true prognostic gene programs in controlled simulations. (A) Test-set concordance index across 100 simulation replicates ($p = 3,000$ genes, $n = 200$ samples, true $k = 3$) comparing DeSurv to standard NMF ($\alpha = 0$); paired Wilcoxon $P < 0.001$. (B) Precision (fraction of a factor’s top-weighted genes, defined in main text, overlapping a true prognostic program) across replicates. In (A) and (B), boxes show median and interquartile range (IQR); whiskers extend to $1.5 \times \text{IQR}$; points beyond are individual replicates. The y-axis in (A) is truncated at 0.5 because C-index values below 0.5 indicate worse-than-random concordance. (C) Distribution of selected k values across the same 100 replicates for DeSurv versus standard NMF. (D) Gaussian-process-predicted mean cross-validated C-index from Bayesian optimization over factorization rank (k) and supervision strength (α) in PDAC training data (TCGA + CPTAC), showing the broad plateau spanning $k = 3$ –12 from which the one-standard-error rule selects $k = 3$ (Methods).